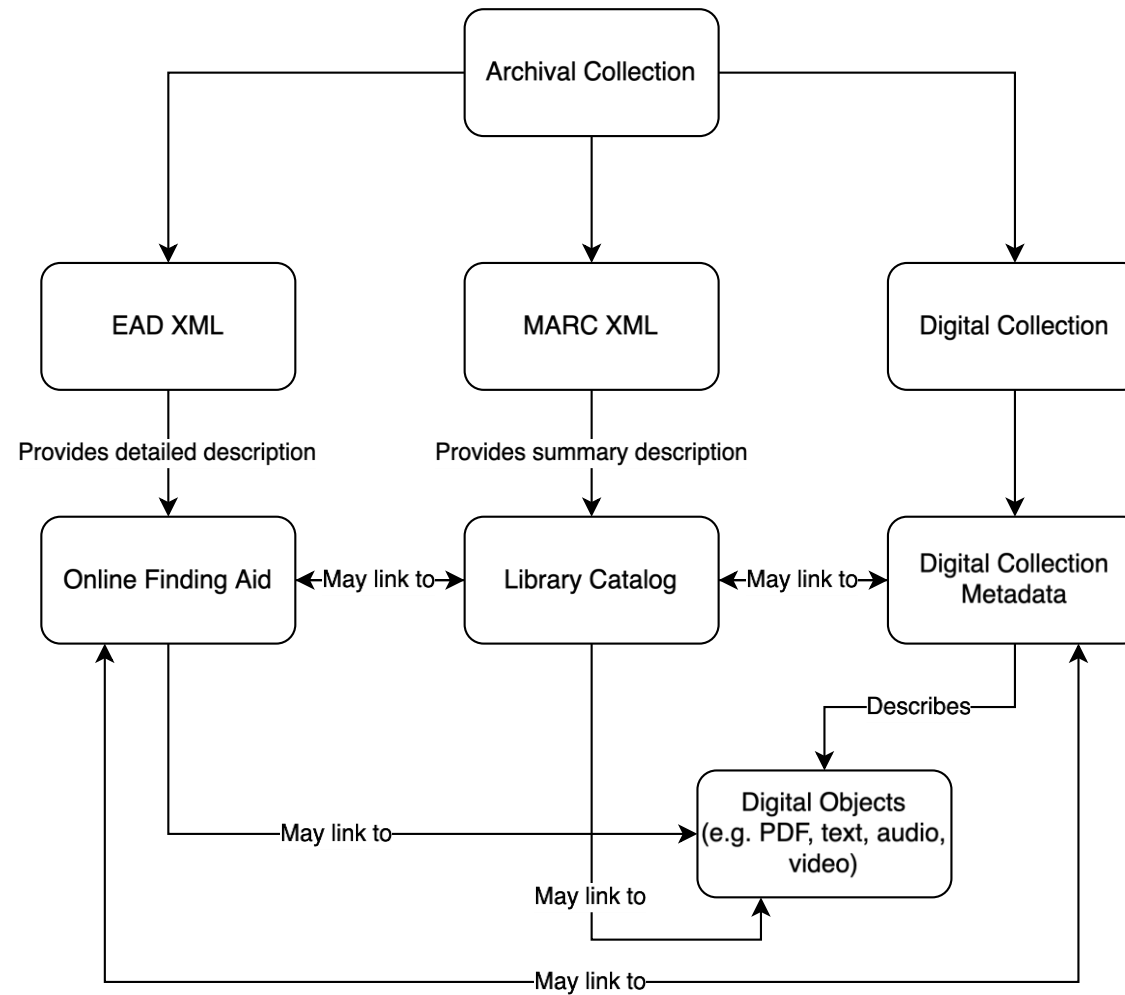


# FolkRAG

A Retrieval-Augmented Generation System for Cultural  
Heritage Materials

# Archival Data

- Collections described across multiple systems:
  - Catalog records (bibliographic)
  - Finding aids (hierarchical)
  - Digital objects (file-level)
- Current limitations:
  - Multiple disconnected systems
  - Complex navigation
  - Requires expert knowledge



# American Folklife Center

- Division at Library of Congress
- Founded 1976 by American Folklife Preservation Act
- Houses Archive of Folk Culture (est. 1928)
- Materials: Folk songs, oral histories, field notes, manuscripts, photographs, audio/video recordings

# FolkRAG

Simple UI that enables researchers and librarians to more efficiently identify their research materials and reference documents.

## Proof of Concept

- Combines LLMs with information retrieval
- Enables natural language queries
- Maintains archival context, provides citations
- 94% retrieval accuracy

# Data Acquisition and Processing

- Web scraping pipeline:
  - EAD XML from LoC finding aids
  - MARC XML from catalog records
  - JSON metadata via LoC API
- Document processing:
  - 10,980 Whisper audio transcriptions
  - 4,292 Tesseract OCR outputs
  - 29,778 library-created text files
- Metadata integration:
  - Pattern matching for identifiers
  - Hierarchical field mapping

# Initial Steps

- Started with FAISS:
  - Fast similarity search
  - GPU optimization
  - Metadata limitations
- Switched to DeepLake:
  - Individual tensor per field
  - Preserved metadata hierarchy
  - Efficient filtering capability

# Technical Overview

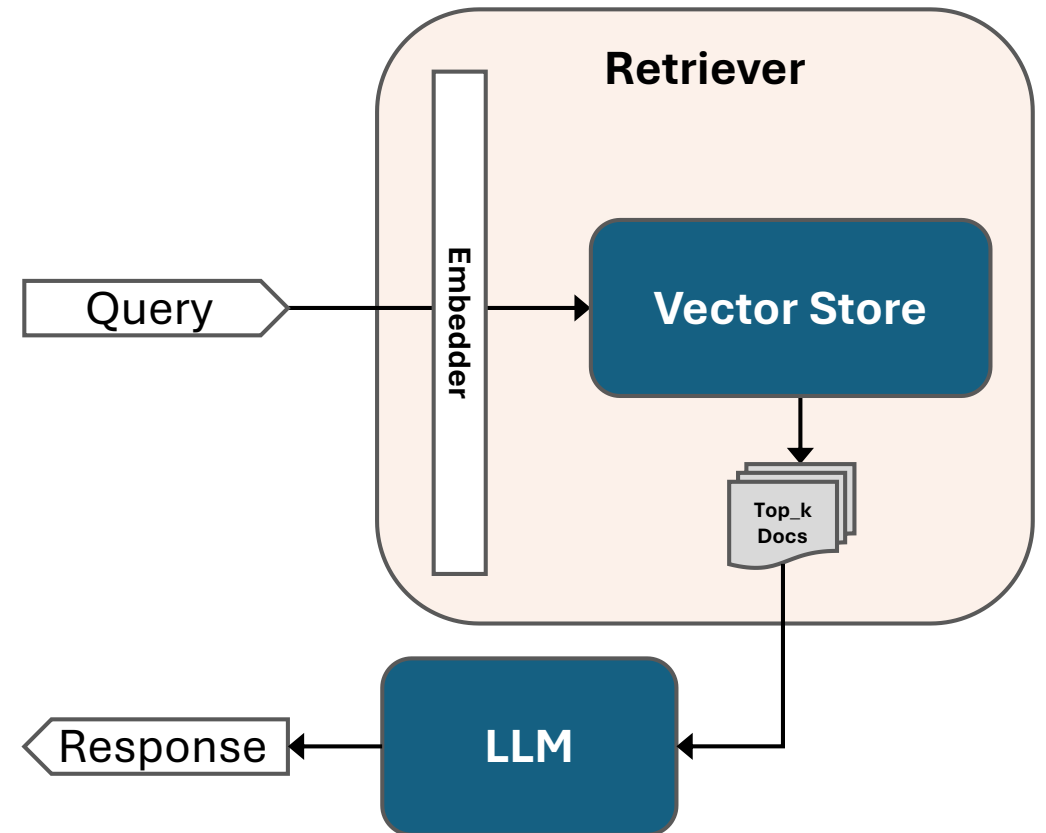
- DeepLake vector store w/metadata
- Instructor-XL embeddings
- HyDE generator (Claude/Bedrock)
- TF-IDF re-ranking



# RAG Architectures: Naïve

## Primary Components

- Vector database
- LLM

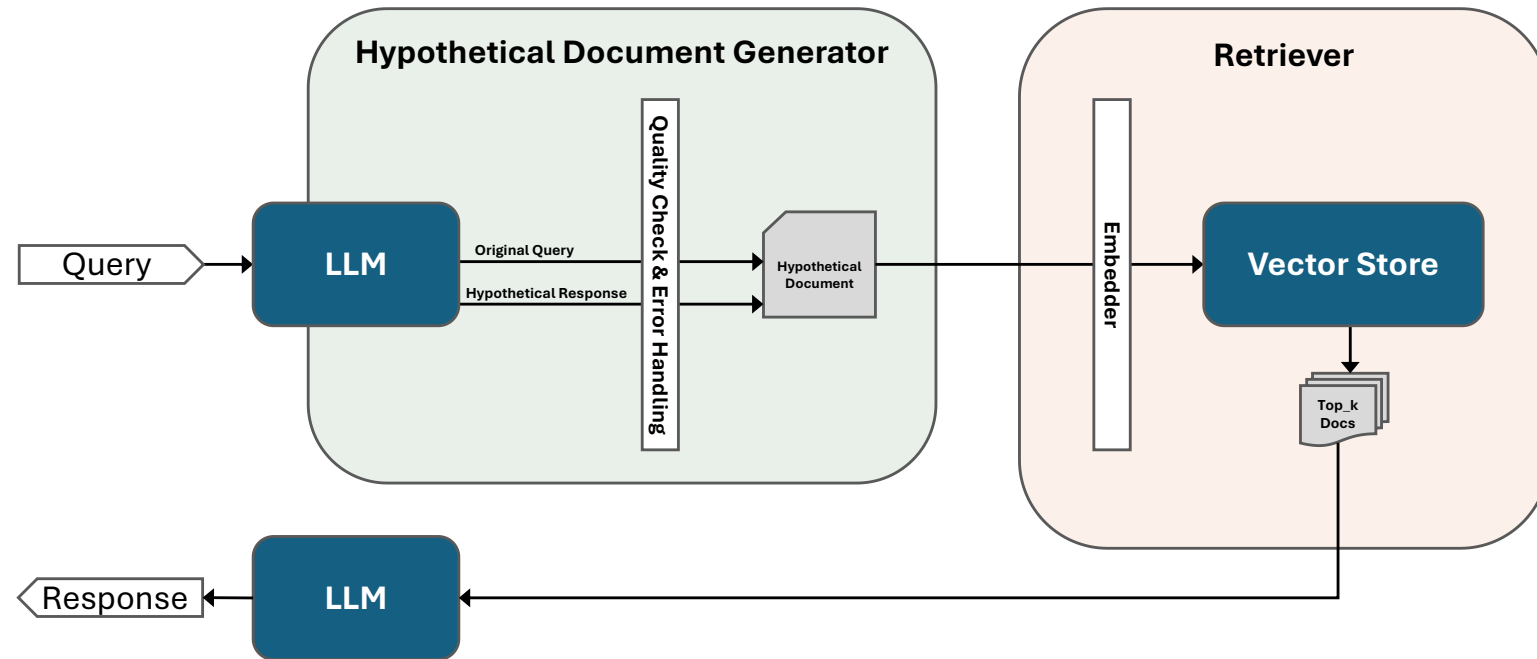


# RAG Architectures: HyDE

Naïve + HyDE

## Primary Components

- HyDE Generator
- Vector database
- LLM



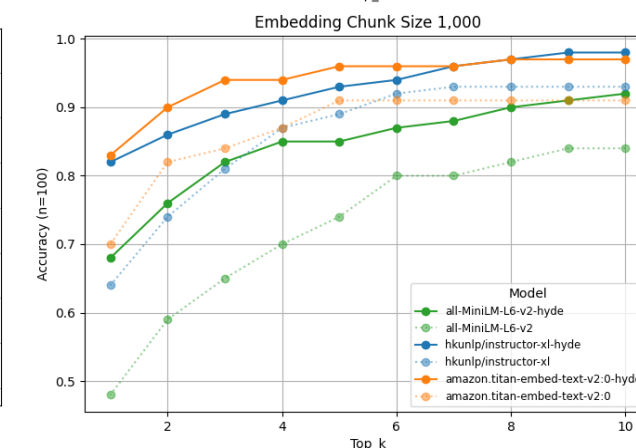
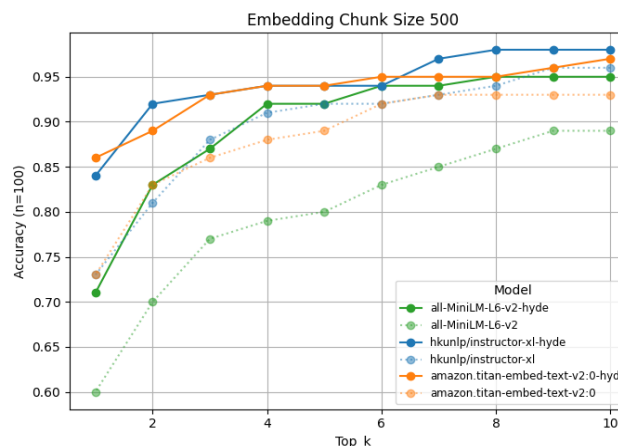
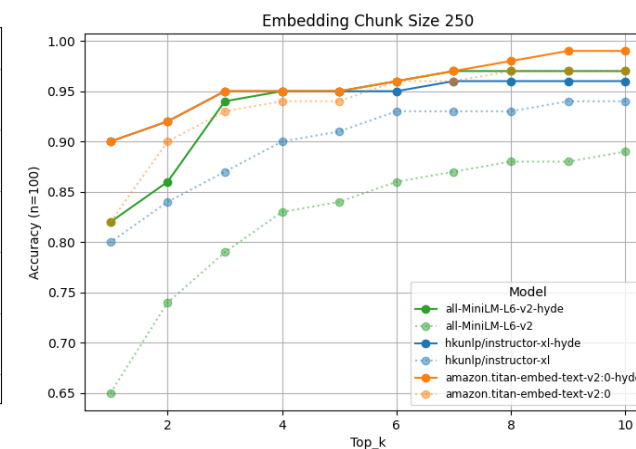
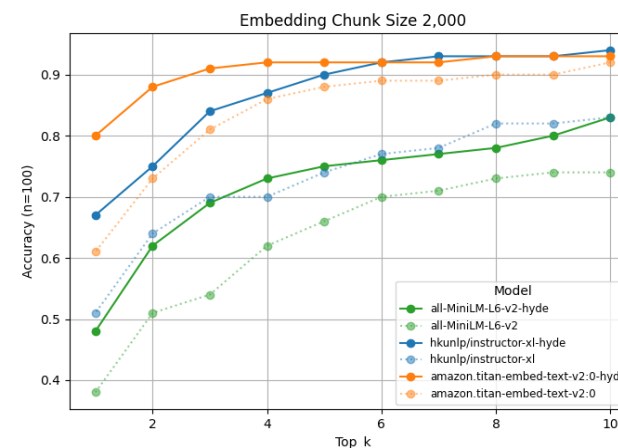
# Evaluation: Sample Data

## Criteria

- Accuracy on 100 test questions (Correct document filename among top\_k documents)
- RAG System (Naïve, HyDE)

## Observations

- Drop MiniLM as an embedding option
- Retain 250, 1,000 sized chunking strategies

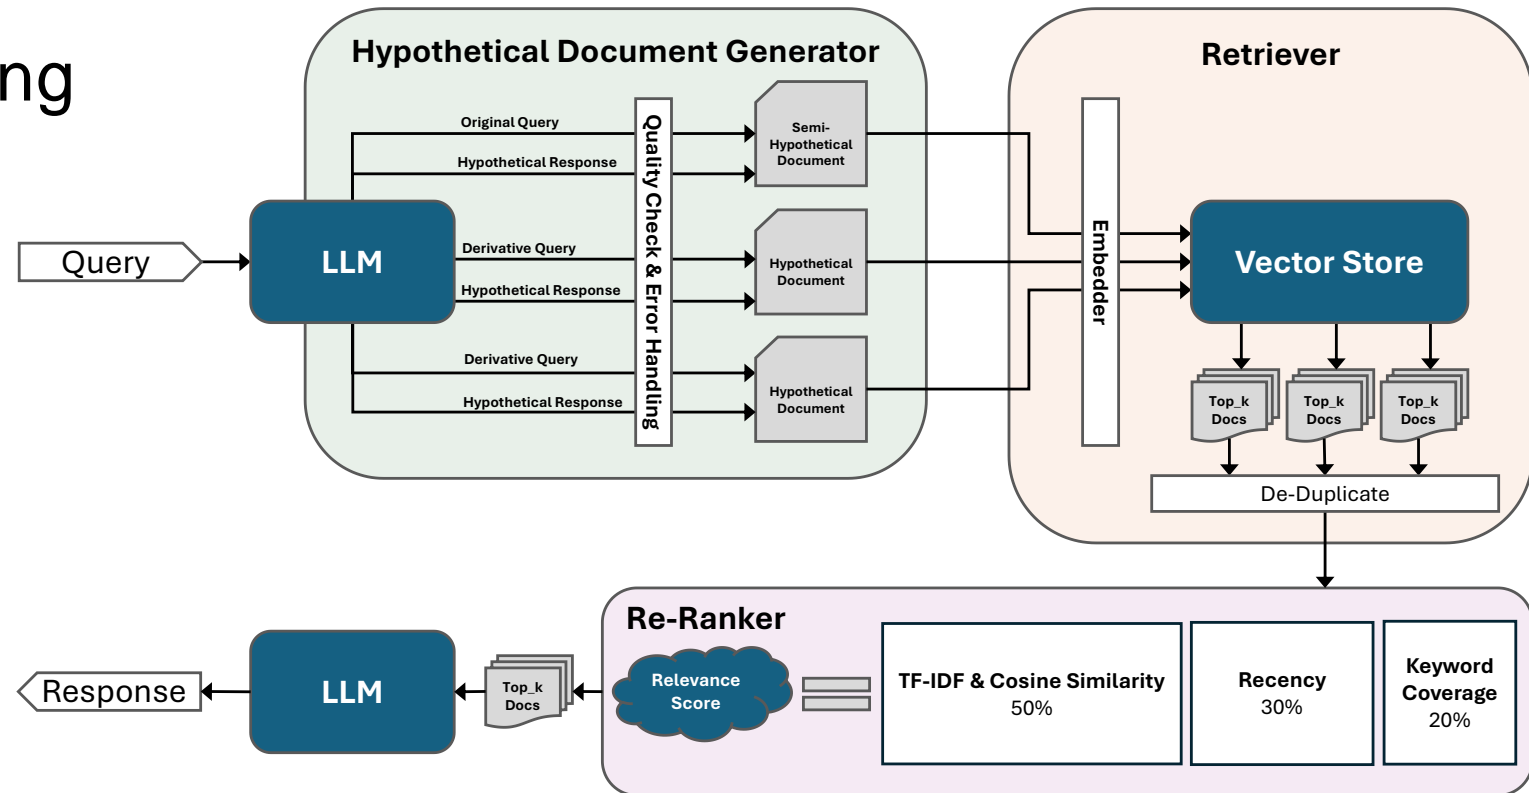


# RAG Architectures: HyDER

Naïve + HyDE + Re-ranking

## Primary Components

- Enhanced HyDE Generator
- Vector database
- Custom Re-ranker
- LLM



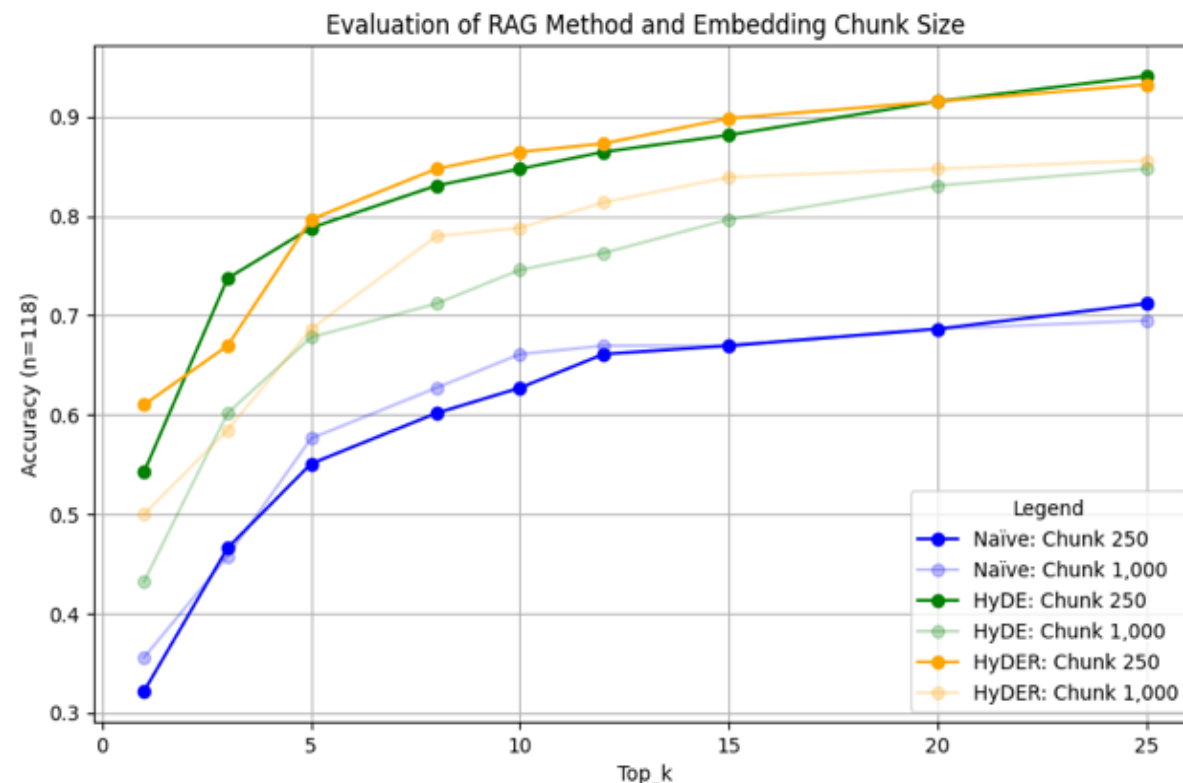
# Evaluation: All Data

## Criteria

- Accuracy on 18 test questions (Correct document filename among top\_k documents)
- RAG System (Naïve, HyDE, HyDER)

## Best Model Config

- Instructor-xl
- Smaller chunking strategy (250)
- HyDER (slight preference)



# Evaluation: Generation

Query: *What is Captain Pearl R. Nye's favorite ballad?*

HyDER (250 Chunked Vector DB)

*Captain Pearl R. Nye's favorite ballad is Barbara Allen. This information can be found in the context, where Captain Nye states, "I guess Barbara Allen is my favorite of them all." The recording was made on June 27, 1937 in Akron, Ohio, as part of a field interview conducted by John A. Lomax. This information is part of a sound recording titled "Two Sailors: Sea Shanties and Canal Boat Ballads," which is one of 10 radio programs in "The Ballad Hunter" radio series. The original filename for this information is **sr28a\_en.txt**.*

HyDER (1,000 Chunked Vector DB)

*Captain Pearl R. Nye's favorite ballad is Barbara Allen. This information can be found in the file **"afc1933001\_ms403\_002.txt"**. The context explicitly states: "I guess Barbara Allen is my favorite of them all." This quote is attributed to Captain Nye in the passage. Additionally, the context mentions that the Library of Congress has eighty-three versions of Barbara Allen, further emphasizing the popularity and significance of this ballad.*

*The metadata indicates that this information comes from the John A. Lomax and Alan Lomax papers, specifically from a manuscript titled "Adventures of a Ballad Hunter." The document was created between 1907 and 1969 in the United States, though Captain Nye's exact recording date is not specified in the given information.*

# Results

- Naive RAG baseline: ~45% accuracy
- Basic HyDE: ~75% accuracy
- HyDER: ~90% accuracy at top-k=15
- Chunk size comparison:
  - 250 chars: 94% max accuracy
  - 1000 chars: 87% max accuracy
  - 2000 chars: 82% max accuracy

# Conclusions and Future Work

- Future Work:
  - Multimodal Integration:
    - Image understanding
    - Cross-modal relationships
    - Visual context preservation
  - Conversational Interface:
    - Memory of past interactions
    - Context accumulation
    - Thread-based exploration